

01 ClearList / AML Alert-Triage

Agent

An LLM judge that auto-clears low-risk transaction-monitoring alerts so investigators only see likely-true cases.

WHY IT WAS BUILT

Meridian's transaction-monitoring system fires roughly 50,000 alerts a month, and the overwhelming majority are false positives. Investigators are over budget and behind on SAR deadlines. The pitch: let the agent clear the obvious noise so humans spend their hours on real candidates.

HOW IT WORKS

- Rules engine generates an alert; the agent reads the case context (customer, transaction history, prior dispositions).
- An LLM judge scores each alert **escalate** vs. **clear**.
- Cleared alerts are auto-dispositioned and closed; escalated alerts go to the human investigator queue.

BY THE NUMBERS

~50,000

alerts / month

~2%

are truly suspicious (base rate)

90% / 5%

judge's claimed TPR / FPR

Days

SAR filing deadline pressure

YOUR TABLE

BSA/AML Officer

Line investigators

Model validator

FinCEN examiner

YOUR HOME TENSION

The Eval Gap

At a 2% base rate, what is the PPV of an “escalate” — and how often is a “clear” actually a missed SAR? Compute the effective sensitivity. Remember: zero flags is not zero failures.

AT THE AI RISK COMMITTEE (CAPSTONE)

Defend auto-clearing to the BSA/AML officer. What is the agent's miss rate at this prevalence, and what audit stream proves it?

02 MemoWriter / Credit-Memo Drafting

Agent

Drafts commercial credit memos from financial spreads and borrower packages so underwriters edit rather than write.

WHY IT WAS BUILT

A junior analyst takes two to three days to draft a single commercial credit memo, and underwriting throughput is the decisioning bottleneck. The pitch: the agent drafts the narrative — risk-rating rationale, covenants, mitigants — and the underwriter reviews instead of writing from scratch.

HOW IT WORKS

- Ingests financial spreads and the borrower document package.
- Drafts the full memo narrative: rating rationale, proposed covenants, identified mitigants.
- Underwriter reviews and edits; credit officer approves the final rating.

BY THE NUMBERS

~3,000

credit memos / year

~18 mo

to train an analyst to independent judgment

Drafting

the routine task being automated

The rating call

the judgment that actually binds

YOUR TABLE

Chief Credit Officer

Senior underwriters

Junior analysts

Tech – AI Platform

YOUR HOME TENSION

Bainbridge's Debt / Weak Links

Drafting is the practice substrate where juniors learned the judgment that limits output quality. If you automate the substrate, where does the next generation of underwriters acquire the binding-constraint skill?

AT THE AI RISK COMMITTEE (CAPSTONE)

Show the Chief Credit Officer your skill-atrophy signal. What did throughput buy this year, and what will it have cost in eighteen months?

03 Navigator / Customer-Service GUI

Agent

| A GUI agent that operates internal servicing systems to resolve customer requests end-to-end.

WHY IT WAS BUILT

The call center is swamped; handle times and after-hours coverage are the pain. The pitch: a GUI agent reads the customer's request and operates the internal teller and servicing screens directly — disputes, address changes, fee reversals, even account actions — to resolve the issue without a human in the chair.

HOW IT WORKS

- Customer chat arrives; the agent interprets the request.
- The agent navigates internal servicing GUIs via screenshots and actions.
- It executes the resolution — some of which are irreversible (move funds, close an account, reverse a fee).

BY THE NUMBERS

~200,000
contacts / month

Seconds
action-to-commit latency

Irreversible
a meaningful slice of actions

Unclear
current human override path

YOUR TABLE

Head of Servicing

Fraud ops

Customers

Legal (Reg E)

YOUR HOME TENSION

GUI Agents + Human and the Loop

Which actions are reversible and which commit immediately? For the irreversible ones, does a human checkpoint run *before* commit — and at a speed and transparency where judgment can actually matter?

AT THE AI RISK COMMITTEE (CAPSTONE)

Bring the servicing head a checkpoint design: which actions gate on a human, at what threshold, and how you avoid an alert-fatigue rubber stamp.

04 SettleNet / A2A Payments Reconciliation

Workflow

A multi-agent workflow that reconciles and posts payments straight-through across internal systems and a vendor.

WHY IT WAS BUILT

Manual reconciliation breaks down at volume and straight-through-processing rates are the target. The pitch: an orchestrator coordinates an internal reconciliation agent and an external vendor's agent over an agent-to-agent protocol, with a high-privilege agent authorized to post entries.

HOW IT WORKS

- An orchestrator agent coordinates low-privilege query agents and a high-privilege posting agent.
- The vendor's agent sits **outside** Meridian's trust boundary and exchanges messages over A2A.
- The posting agent can move money across wire, RTP, and ACH — each with different liability rules.

BY THE NUMBERS

\$ billions

flowing per day

Outside

vendor agent vs. trust boundary

High-privilege

the posting agent can move money

3 rails

wire / RTP / ACH, different liability

YOUR TABLE

Treasury ops

Vendor risk

Tech – Infosec

Legal

YOUR HOME TENSION

Agent to Agent + Adversarial Testing

If the low-privilege vendor agent is compromised, what stops a cross-agent injection from pivoting into the high-privilege posting agent? And when a wrong payment commits, which rail's rules decide who eats it?

AT THE AI RISK COMMITTEE (CAPSTONE)

Draw the trust boundary for Treasury and infosec. Show the control that prevents a low-privilege compromise from reaching the posting agent.

05 Second Line / Model-Validation

Copilot

An agent that helps MRM validators review documentation, reproduce results, and draft validation findings.

WHY IT WAS BUILT

The validation backlog is growing faster than the bank can hire validators, and SR-cycle pressure is real. The pitch: the copilot checks documentation completeness, reproduces reported metrics, and drafts findings so scarce validators move faster.

HOW IT WORKS

- Ingests model documentation and code.
- Checks documentation completeness, reproduces reported metrics, drafts validation findings.
- A human validator reviews — trusting, in practice, the copilot's “this looks complete.”

BY THE NUMBERS

~400

models in inventory

Growing

validation backlog

Footnote 3

SR 26-2 excludes GenAI from scope

An LLM

the copilot's judgment is itself a model

YOUR TABLE

Head of MRM

Model validator

Model owners

Internal audit

YOUR HOME TENSION

Regulatory Drift + Effective Challenge

The guardrail is itself an LLM that needs validation — who validates the validator? And if SR 26-2 footnote 3 excludes GenAI from scope, is this copilot even “a model” under the framework you'd use to govern it?

AT THE AI RISK COMMITTEE (CAPSTONE)

Tell the Head of MRM whether the copilot is in scope, who provides effective challenge to it, and what its false-negative rate on findings is.

△ A curveball envelope will reach your pod mid-course — do not open until told.

Pod 05 · Second Line

06 Compass / Wealth-Advisory Research

Agent

Drafts investment research and portfolio recommendations for advisors at scale.

WHY IT WAS BUILT

Research capacity per advisor is the constraint and cost-per-note is the headline metric. The pitch: the agent turns market data and house views into research notes and suggested allocations, so 1,500 advisors can each serve more clients.

HOW IT WORKS

- Combines market data with the firm's house views.
- Drafts research notes and suggested client allocations.
- The advisor reviews before the recommendation reaches the client.

BY THE NUMBERS

1,500

advisors served

Up to 420×

output-token price spread across model tiers

Activity vs. outcome

which is the bank measuring?

Drift

toward higher-risk recommendations

YOUR TABLE

Head of Wealth

Advisors

Tech – AI Platform

Compliance (Reg BI)

CFO

YOUR HOME TENSION

Economics of AI

Are you measuring cost-per-outcome or cost-per-activity? As advisors lean on the agent, who still forms an independent view — and what catches a slow drift toward higher-risk recommendations before a regulator does?

AT THE AI RISK COMMITTEE (CAPSTONE)

Give the CFO and compliance the unit economics and the drift-detection cadence. Cost-per-outcome, not notes-per-day.

△ A curveball envelope will reach your pod mid-course — do not open until told.

Pod 06 · Compass