

# Two-Day Course — Facilitator Schedule

10 lecture-modules over two days · up to 6 pods of 3–5 · on paper & whiteboard (no sandbox required)

**The frame.** The whole course runs inside one fictional bank, **Meridian National**, standing up its GenAI portfolio. Each pod owns one live use case (its case brief) for both days and defends it at a mock AI Risk Committee in the capstone. Day 1 is **build your case**; Day 2 is **defend it under fire** as curveballs land.

## THE SPINE – LECTURE VS. STUDIO

Eight modules are short lectures feeding a team exercise. Two — **The Eval Gap** and **Adversarial Testing** — are run as **studios**, taught entirely through the exercise (the base-rate computation and the red-team relay), because for this audience doing beats hearing. Modules 7+8 and 9+10 are paired into single slots to fit two days.

## RECURRING MECHANICS

■ Lecture ■ Studio ■ Exercise ■ Mechanic ■ Logistics

Cold-open (vote before the lecture) · steelman-swap (argue the opposite) · module quiz as a team leaderboard · gallery walk as the mock committee. None require devices.

## POD → CASE → CURVEBALL MAP

| POD | CODENAME    | USE CASE                 | HOME TENSION                 | CURVEBALL LENS      |
|-----|-------------|--------------------------|------------------------------|---------------------|
| 01  | ClearList   | AML alert-triage         | The Eval Gap                 | Human and the Loop  |
| 02  | MemoWriter  | Credit-memo drafting     | Bainbridge's Debt            | Economics of AI     |
| 03  | Navigator   | Customer-service GUI     | GUI Agents + HITL            | Adversarial Testing |
| 04  | SettleNet   | A2A payments recon       | Agent to Agent               | Regulatory Drift    |
| 05  | Second Line | Model-validation copilot | Regulatory Drift + challenge | The Eval Gap        |
| 06  | Compass     | Wealth-advisory research | Economics of AI              | Bainbridge's Debt   |

All six curveballs open together at the end of Day 1; pods carry the twist overnight and arrive Day 2 ready to develop a response. Each curveball deliberately targets a lens that is *not* the pod's home tension, so every pod has to stretch into a second domain.

## Day 1 — Build your case

| TIME  | BLOCK                                      | TYPE      | FACILITATOR NOTE                                                                                                                                               |
|-------|--------------------------------------------|-----------|----------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 9:00  | Welcome, pod formation & case draw         | LOGISTICS | Form up to 6 pods of 3–5. Each pod draws one Meridian National case brief and keeps it for two days.                                                           |
| 9:00  | Cold-open provocation                      | MECHANIC  | Three motivations for GenAI (throughput / cost / “everyone’s doing it”). Pods vote the strongest, then learn each case was greenlit for one of them.           |
| 9:30  | M1 · Regulatory Drift                      | LECTURE   | SR 11-7 → SR 26-2; the GenAI scope exclusion (fn 3).                                                                                                           |
| 10:20 | Exercise A · Scope triage                  | EXERCISE  | Pods classify 5 use cases as in / out of SR 26-2 scope and defend each call on the whiteboard.                                                                 |
| 10:45 | Break                                      | LOGISTICS |                                                                                                                                                                |
| 11:00 | M2 · Agentic Engineering                   | LECTURE   | Harness design, the agent loop, where control lives. Pod 02’s home module.                                                                                     |
| 11:50 | Exercise B · Harness teardown              | EXERCISE  | Each pod is handed a flawed agent harness sketch and finds the failure mode.                                                                                   |
| 12:15 | Lunch                                      | LOGISTICS |                                                                                                                                                                |
| 13:15 | M3 · Agent to Agent                        | LECTURE   | Orchestration, trust boundaries, privilege. Pod 04’s home module.                                                                                              |
| 14:05 | Exercise C · A2A threat model              | EXERCISE  | Pods map the trust boundary and privilege levels on their own case.                                                                                            |
| 14:30 | Break                                      | LOGISTICS |                                                                                                                                                                |
| 14:45 | M4 · GUI Agents                            | LECTURE   | Screen-driving agents, irreversible actions. Pod 03’s home module.                                                                                             |
| 15:35 | Project work block                         | EXERCISE  | Pods assemble their Day-1 “build” artifact on a flip chart — architecture + where the risk lives.                                                              |
| 16:05 | Module quiz competition                    | MECHANIC  | Run the existing module tests as a team leaderboard game (no devices needed; read aloud or print).                                                             |
| 16:30 | Curveball reveal — all pods open envelopes | MECHANIC  | Every pod opens its sealed envelope at the same time. First reactions only — pods carry the twist overnight and develop their response for the Day-2 capstone. |
| 16:50 | Day 1 close                                | LOGISTICS | Build your case today. Sit with the curveball tonight. Defend it tomorrow.                                                                                     |

All six curveballs open together at the Day-1 reveal block. Adjust block lengths to your room; the exercises compress to 15 min or stretch to 35 without breaking the arc.

## Day 2 — Defend it under fire

| TIME  | BLOCK                                            | TYPE      | FACILITATOR NOTE                                                                                                                                       |
|-------|--------------------------------------------------|-----------|--------------------------------------------------------------------------------------------------------------------------------------------------------|
| 9:00  | Recap + curveball regroup                        | MECHANIC  | Quick glossary buzzer round to warm the room; pods take five minutes to compare the overnight thinking on their curveball.                             |
| 9:20  | M5 · The Eval Gap — STUDIO                       | STUDIO    | Taught through the exercise. Pod 01's home module.                                                                                                     |
| 9:20  | Studio · The base-rate trap                      | EXERCISE  | Pods compute PPV for an LLM judge at low prevalence and discover zero flags is not zero failures.                                                      |
| 10:10 | Steelman-swap                                    | MECHANIC  | Each pod argues the case <i>against</i> its own instinct: "the judge is good enough."                                                                  |
| 10:40 | Break                                            | LOGISTICS |                                                                                                                                                        |
| 10:55 | M6 · Adversarial Testing — STUDIO                | STUDIO    | Taught through the relay.                                                                                                                              |
| 11:25 | Studio · Red-team relay                          | EXERCISE  | Each pod writes 3 attacks, passes them clockwise; the receiving pod defends.                                                                           |
| 12:10 | Lunch                                            | LOGISTICS |                                                                                                                                                        |
| 13:10 | M7 + M8 · Human and the Loop / Bainbridge's Debt | LECTURE   | Paired lecture. Pods 01 and 06 deepen their curveball response here.                                                                                   |
| 14:10 | Exercise · Knight Capital / override signal      | EXERCISE  | Pods design one runtime oversight signal for their case (override rate, drift cadence, deskilling proxy).                                              |
| 14:35 | Break                                            | LOGISTICS |                                                                                                                                                        |
| 14:50 | M9 + M10 · Banking AI / Economics of AI          | LECTURE   | Paired lecture. Pod 02 deepens its curveball response here.                                                                                            |
| 15:40 | Capstone · Gallery walk (mock AI Risk Committee) | MECHANIC  | Each pod posts its case, build, and how it handled its curveball, then defends a go/no-go. The room rotates and challenges. Effective challenge, live. |
| 16:30 | Final-exam leaderboard + certificate + close     | MECHANIC  | Optional team final exam; award the completion certificate.                                                                                            |

M5 and M6 are studios — the exercise *is* the teaching. M7+M8 and M9+M10 are paired lectures. Pods weave their already-open curveball into each relevant module, then present the result at the capstone.

## Exercise playbook (1 / 2)

### Cold-open provocation

Day 1 · 9:00 · 15 min

Open with a single contrarian claim on a slide; pods vote agree/disagree on the whiteboard before any lecture. Reveal the tension the two days will test. Reuse the format before each module as a 60-second warm-up.

Output: a visible split in the room you can call back to later.

### Exercise A · Scope triage

Day 1 · M1 · 25 min

Hand each pod five short use-case descriptions. They sort each in / out of SR 26-2 scope and write the one sentence that decides it. Surface the GenAI exclusion as the trap.

Output: a sorted list + the deciding sentence per case.

### Exercise B · Harness teardown

Day 1 · M2 · 25 min

Give every pod the same one-page flawed harness sketch (e.g. no termination contract, unbounded tool access). Pods mark the failure mode and the smallest fix.

Output: one circled failure + one-line fix.

### Exercise C · A2A threat model

Day 1 · M3 · 25 min

Pods draw their own case as agents + privileges on the whiteboard, then mark the trust boundary and the highest-privilege action.

Output: a boundary line and the one action that needs the most scrutiny.

### Studio · The base-rate trap

Day 2 · M5 · 50 min (taught as the lecture)

Give the prevalence and the judge's TPR/FPR. Pods compute PPV by hand, then the effective sensitivity (judge × reviewer). The math is the lecture. ClearList's numbers are on its brief.

Output: a PPV figure and the sentence "zero flags ≠ zero failures" in their own words.

### Studio · Red-team relay

Day 2 · M6 · 60 min (taught as the lecture)

Round 1: each pod writes 3 attacks against its own case. Round 2: pass clockwise; the receiving pod classifies (OWASP/ASI) and designs a control. Round 3: brief back.

Output: 3 attacks + 3 controls per case, classified.

## Exercise playbook (2 / 2)

---

### Exercise • Knight Capital / override signal

Day 2 • M7+M8 • 25 min

Reenact the 97-emails failure briefly, then each pod designs one runtime oversight signal for its case: override rate, drift cadence, or a deskilling proxy.

Output: one signal, its threshold, and who watches it.

### Steelman-swap

Recurring • ~10 min

A pod must argue the position opposite to its instinct (e.g. “runtime monitoring is enough”). Forces honest engagement with the counterargument the course keeps raising.

Output: the strongest version of the view they disagree with.

### Module quiz competition

Day 1 close • 30 min

Run the blog's existing module tests as a team game — read questions aloud or print them. Leaderboard on the whiteboard. No devices required.

Output: a leaderboard going into Day 2.

### Capstone • Gallery walk

Day 2 • 15:40 • 50 min

Each pod posts its case, its build, the curveball it absorbed, and its go/no-go recommendation. The room rotates as a mock AI Risk Committee and challenges each. This is effective challenge, performed.

Output: a committee verdict per case + the strongest challenge each pod received.

**Materials checklist (per pod):** case brief (printed), whiteboard or flip chart + markers, the curveball envelope (sealed), printed module-quiz questions, and one calculator or phone for the base-rate studio only.