

A TWO-PART SERIES · APRIL 2026

SR 11-7 at a Crossroads

Model risk management from the 2008 crisis to GenAI — and which 2011 idea still has to survive the transition.

PART 1

SR 11-7 at Fifteen

PART 2

Effective Challenge Is Worth Keeping

Two parts, one argument

Part 1 lays out why the 2011 framework still governs and where GenAI breaks it. Part 2 isolates the one piece — effective challenge — that survives, and asks what it would take to preserve.

PART 1 March 25, 2026

SR 11-7 at Fifteen

A 2011 framework designed for regression models now governs GenAI. Regulators have signaled change but not delivered it.

01 Crisis-era origins

02 Three pillars + effective challenge

PART 2 April 25, 2026

Effective Challenge Is Worth Keeping

The triad — incentives, competence, influence — is the only 2011 insight that survives. Each leg is now under acute stress.

PART 1 · MARCH 25, 2026

SR 11-7 at Fifteen

Model risk management faces an AI reckoning.

The failures that made SR 11-7 inevitable

1998 • LTCM

VaR models assumed constant volatility and stable correlations. Calibration to a benign period catastrophically understated tail risk.

2000 • Gaussian copula

David Li's model assumed default correlations were constant. Modeled at 0.2–0.3, they surged to nearly 1.0 when housing turned.

2008 • AIG

CDS protection sold at near-zero premiums on the same flawed correlation assumption. The bailout cost \$182 billion.

EMANUEL DERMAN, 1996

"There are always implicit assumptions behind a model and its solution method. But human beings have limited foresight and great imagination, so that, inevitably, a model will be used in ways its creator never intended."

Three pillars and one definition

01

Sound development

Clear purpose, rigorous methodology, documentation an unfamiliar party can follow.

02

Rigorous validation

Conceptual soundness, ongoing monitoring, outcomes analysis — backtesting, sensitivity, benchmarking.

03

Governance

Board-level oversight, senior management accountability, full model inventory.

EFFECTIVE CHALLENGE

"Critical analysis by objective, informed parties who can identify model limitations and assumptions and produce appropriate changes."

THE PHRASE THAT BECOMES LOAD-BEARING

A model produces "**quantitative estimates**" — wording that becomes a definitional battleground fifteen years later when the output is a paragraph, not a number.

What the first decade revealed

2012	JPMorgan's London Whale A modified VaR model dramatically understated risk on a synthetic credit derivatives portfolio. Internal controls did not detect it.	\$6.2B trading loss · \$300M OCC civil money penalty in 2013
2020	Current Expected Credit Losses (CECL) meets COVID Forward-looking lifetime loss estimates were barely implemented when the pandemic overwhelmed them. Management overlays filled the gap.	82% of banks: credit-risk models most affected · 66% used overlays
—	The vendor model compromise SR 11-7 required vendor models to enter the same MRM framework — but banks could not replicate or inspect proprietary methodology.	A compromise in 2011. Fatal when the vendor is OpenAI or Anthropic.

GenAI breaks each core assumption

CONCEPTUAL SOUNDNESS

SR 11-7 assumed

~~A model with reviewable parameters and a traceable derivation.~~

GenAI delivers

A network with billions of parameters and no clear analog of 'conceptual soundness'.

Validation lacks the surface area it was designed to operate on.

REPRODUCIBILITY

SR 11-7 assumed

~~Same inputs produce the same outputs each time.~~

GenAI delivers

The same prompt produces different outputs. Wording changes alter behavior.

Backtesting and outcomes analysis assume an invariance that no longer holds.

PERIODIC REVIEW

SR 11-7 assumed

~~A stable model form, reviewed on a cycle the bank controls.~~

GenAI delivers

Multi-step agentic systems that take autonomous action between reviews.

The review cadence is shorter than the model's window of unmonitored action.

A patchwork — then a carve-out

United States · fragmented

- OCC** **Bulletin 2025-26** · A 'first step' in a broader MRM review
Acknowledges over-application to low-risk models at smaller banks.
- FED** **Governor Barr** · April 2025 speech
'Should review and update existing standards' — without the specificity of formal guidance.
- TRES.** **Treasury AI RMF** · February 2026
230 control objectives. Voluntary, but expected to shape examiner expectations.

International · moving faster

- UK** **PRA SS1/23** · May 2023
Adopts 'effective challenge' verbatim. Treats model risk as a risk in its own right.
- CA** **OSFI E-23** · Effective May 2027
Technology-neutral definition explicitly covering AI/ML. The closest existing template for a rewrite.
- EU** **AI Act** · Credit scoring = high-risk
Binding documentation, explainability, oversight. Compliance deadline August 2, 2026.

Apr 17, 2026

SR 11-7 is superseded by joint Fed/OCC/FDIC SR 26-2. GenAI and agentic AI are explicitly excluded from scope.

PART 2 · APRIL 25, 2026

Effective Challenge Is Worth Keeping

The rest is negotiable.

The triad rests on three legs

01 Incentives

Separate development, supported by culture and incentives — so the validator is not measured the way the builder is.

FAILURE MODE

Structural separation that doesn't survive a pay gap.

02 Competence

Technical knowledge to evaluate — to understand 'how and why' the model produces what it does.

FAILURE MODE

A validator who cannot trace effects to causes.

03 Influence

Organizational authority over deployment, backed by senior management commitment that defends second-line authority.

FAILURE MODE

Authority on paper that no one is willing to invoke.

OCC, 2021 Handbook: "Effective challenge is synonymous with credible challenge" — the shift from "does an independent team exist?" to "is the challenge producing risk reduction?"

The failure modes GenAI amplifies

01 Incentives

BIS pay-and-experience gap

First-line builders are better-paid and more senior than the validators challenging them.

Vendor feature removal

Banks ask vendors to disable AI features rather than survive an MRM review.

Technology-org gatekeeping

GenAI arrives via SDLC governance, not MRM. Ownership is unresolved.

02 Competence

FINMA, 2023

"Some AI model results cannot be understood, explained, or reproduced — and therefore cannot be critically assessed."

Non-determinism

Identical inputs yield different outputs. Backtesting assumes invariance that no longer holds.

Test-aware models

LLMs can alter their responses when they detect they are being evaluated.

03 Influence

Vendor release cadence

A bank cannot "retire" an LLM when the vendor updates the weights on its own schedule.

Agentic ownership lines

Pipelines cross approvals. Each component is approved; the chain is not.

C-suite deployment mandates

Top-down adoption velocity signals that second-line authority is subordinate.

The London Whale was never a competence failure

\$6.2B

Trading loss from a new VaR model that cut reported risk by ~50%.

300+

VaR limit breaches resolved by swapping in a more lenient model.

SUM

An Excel error used SUM instead of AVERAGE in the volatility calc.

RISK MANAGERS, POST-MORTEM

"Not sufficiently independent to challenge the office's practices."

Structural separation was there. The **incentives leg** was not. The post-mortem reads like an influence failure with a technical veneer — exactly the failure mode SR 11-7 was written to prevent.

Adaptations on the table

01 Red-teaming as formal challenge

Adversarial testing by a structurally independent team.

Risk: Defaults to the same team probing guardrails it already knows.

02 Risk-tiered governance

Control intensity calibrated to materiality and use case.

Risk: Tier definitions become the new compliance theater.

03 Champion-challenger as default

Alternative models run in parallel with production.

Risk: Both models maintained by the same team — no actual challenge.

04 Continuous monitoring

From periodic validation to ongoing measurement.

Risk: A dashboard generating signals nobody has standing to escalate.

05 Kill switches with disconnection metrics

Pre-established thresholds for autonomous-system shutoff.

Risk: The disconnection metric itself needs its own challenge.

THE COMMON RISK

Each adaptation reproduces the same calibration problem — apparatus standing in for principle — in a new medium.

What holds, what's at risk

01 The triad is load-bearing

Asymmetric incentives, inadequate competence, absent influence — exactly the failures GenAI sharpens. Dropping the framework to escape the bureaucracy creates a larger problem than the one it solves.

02 The influence leg is the watch item

Authority to force change on a foundation model means telling a business owner they can't deploy a vendor product the vendor is actively shipping to their competitors. Nothing in the current conversation says where that authority will live.

03 SR 26-2 deferred the question

The April 2026 guidance carved GenAI out of scope. The next forcing function — a high-profile failure, a foreign regulatory move, a quiet examiner action — will supply the pressure that 2026's formal guidance did not.