

A THREE-PART SERIES · APRIL 2026

Google ADK & A2A: Building, Risks, Reality

From framework primitives to confused-deputy failures — the multi-agent stack, the risk surface it opens, and the cases that prove it isn't theoretical.

PART 1

The Framework

PART 2

The Risk Surface

PART 3

Five Failures

Three posts, one through-line

Part 1 explains the framework Google has shipped for the multi-agent problem. Part 2 maps the risk surface it leaves open. Part 3 walks five real failures that show the surface is already being exploited.

PART 1 April 4, 2026

The Framework

ADK gives six primitives and a cooperative event loop. A2A makes them speak across vendor boundaries on plain HTTP.

PART 2 April 6, 2026

The Risk Surface

Multi-agent compounds risk at the seams. Authentication says who you are talking to, not whether they will behave.

01 Authentication $\neq$ intent

PART 3 April 7, 2026

Five Failures

Knight Capital, EchoLeak, Copilot→Claude, ForcedLeak, TeamPCP, peer-preservation — the confused deputy in production.

Google ADK and the A2A Protocol

A technical reference for the multi-agent stack — six primitives, the cooperative event loop, and a protocol that speaks plain HTTP.

Build with ADK, equip with MCP, communicate with A2A

Google's public framing of what an agent system should be built from. Three roles, three open-spec layers, one stack.

CONSTRUCTION

ADK

Code-first framework for stateful agents. Python, async generators, model-agnostic via LiteLLM.

TOOLS & DATA

MCP

Connects agents to APIs, files, search indices, databases. Vertical integration.

PEER COMMUNICATION

A2A

Open standard for cross-framework agent dialogue. HTTP, JSON-RPC, SSE. Linux Foundation governance.

Six primitives carry the whole framework

Every higher-level pattern in ADK is composed from these six. Learn them and the rest of the framework reads as application.

Agent

The execution unit. All agent types inherit from BaseAgent and yield Event objects from an async generator.

Runner

Orchestrator. Manages the event loop, the session/artifact/memory services, and is the entry point for every invocation.

Session

A single conversation thread — id, user, app, a dict-like state scratchpad, and a chronological event list.

Event

The atomic communication unit. Carries author, content, and an actions field with state_delta / artifact_delta / transfer / escalate.

Artifact

Named, versioned binary data. PDFs, images, generated reports. Auto-incrementing version per save_artifact.

Tool

Anything an agent can invoke — Python functions, sub-agents, built-ins, MCP endpoints, or imported OpenAPI specs.

Five agent types, one spectrum

From fully deterministic workflow to fully LLM-driven reasoning. They mix freely in a single hierarchy.

LlmAgent	LLM-DRIVEN	The reasoning unit. Tools, sub-agents, output_key for piping results between steps.
SequentialAgent	DETERMINISTIC	Sub-agents run in strict order. No LLM in the flow control. Data passes through session state.
ParallelAgent	CONCURRENT	All children run at once. Events interleave; same state key from two children is a race.
LoopAgent	ITERATIVE	Repeats until max_iterations or escalate=True signal. The escape valve for unbounded refinement.

Eight steps. One invariant.

State changes are only guaranteed persisted **after the carrying event has been yielded** and processed by the Runner. Writing then immediately reading inside the same step reads stale state.

- 01** User input arrives — `runner.run_async` is called
- 02** Runner appends user message as an Event via `SessionService`
- 03** Runner calls `agent.run_async`, receives an async generator
- 04** Agent yields an Event — execution pauses at the yield point
- 05** Runner commits the event (state & artifact deltas)
- 06** Runner yields the event upstream to caller / parent
- 07** Agent resumes — can now reliably read what it just wrote
- 08** Cycle repeats until the generator is exhausted

Prefix-scoped state, with policy as a feature

State keys are scoped by prefix. **app:** can be written at configuration time and is read-only from a session context — the pattern for policy parameters users can't manipulate.

SESSION

(none)

Persists within the current session only. The default working scratchpad.

`last_query`

APPLICATION

app:

Shared across all sessions and users. Write-restricted to higher-privilege ops.

`app:max_credit_limit`

PER-USER

user:

Across all sessions for a specific user. KYC, preferences, risk tier.

`user:risk_tier`

INVOCATION

temp:

Discarded after the request-response cycle completes. Volatile.

`temp:retry_count`

Return None to observe. Return a value to intercept.

The return-value-as-interceptor pattern makes short-circuiting clean. `before_model_callback` is the right place for guardrail checks — you skip the LLM call entirely on a fail. The guardrails post covers more.

CALLBACK	TRIGGER	RETURN NONE	RETURN VALUE
<code>before_agent_callback</code>	Before agent runs	Agent proceeds	Skips agent, uses Content
<code>after_agent_callback</code>	After agent completes	Original output kept	Replaces output
<code>before_model_callback</code>	Before LLM API call	LLM call proceeds	Skips LLM, uses LlmResponse
<code>after_model_callback</code>	After LLM response	Original response kept	Replaces response
<code>before_tool_callback</code>	Before tool executes	Tool runs normally	Skips tool, uses dict
<code>after_tool_callback</code>	After tool returns	Original result kept	Replaces result

A2A in three pieces

Open protocol for agents that don't share a codebase, framework, or operator. v1.0 in early 2026; Linux Foundation governance with 150+ supporting organizations. Plain HTTP, JSON-RPC 2.0, SSE — nothing proprietary.

01 · Discovery

Agent Cards

JSON document at /.well-known/agent-card.json declaring name, endpoint, capabilities, auth, skills, modalities.

JWS-signed cards
verifiable client-side.

02 · Unit of work

Task lifecycle

submitted → working → completed. Plus input-required, auth-required, canceled, failed, rejected — the states REST APIs usually pretend don't exist.

contextId groups related multi-turn tasks.

03 · Transport

Three modes

Request/response (message/send). Streaming (SSE for long outputs). Push notifications (webhooks for hours-long, disconnected clients).

Same data model across all three modes.

A2A is for peers. MCP is for tools.

MCP handles vertical integration. A2A handles horizontal. A LoanProcessor uses MCP to call a credit-check API; it uses A2A to delegate disbursement to another agent.

DIMENSION	A2A	MCP
Purpose	Agent-to-agent	Agent-to-tool
Relationship	Peer collaboration	Tool invocation
Unit of work	Task (stateful)	Tool call (structured I/O)
Discovery	Agent Cards	Server manifest
Governance	Linux Foundation	Anthropic

In ADK's cooperative event loop, state changes are 'only guaranteed persisted after the carrying event has been yielded and processed by the Runner.' What subtle bug does this prevent?

- A** It prevents agents from overwriting each other's state during simultaneous parallel execution across concurrent sessions.
- B** An agent reading state it just wrote before yielding reads an uncommitted value.
- C** It prevents the agent from accumulating too many uncommitted state changes before flushing them to storage.
- D** It ensures state changes are encrypted before being stored durably by the SessionService at rest.

When Agents Talk to Agents

Multi-agent compounds risk at the seams — trust dissolves across handoffs, prompt injection propagates like a pathogen, and cascading failures travel faster than incident response.

Authentication says who. It says nothing about what.

A deputy with valid credentials and corrupted instructions looks right at every checkpoint. That's what makes it harder to catch than a bad actor with no access.

- **Agent Cards prove identity, not intent**

Adversarial agents can lie about their capabilities — exaggerating skills to route sensitive tasks to impersonating servers (CSA MAESTRO, Feb 2025).

- **Traditional IAM assumed bounded sessions**

It was designed for human users with predictable action sequences. Autonomous agents fire hundreds of actions within a single authenticated session.

- **Identity validation tells you almost nothing**

Validating at session start says nothing about what the agent actually does for the next forty minutes.

Injection becomes contagion

The injection mechanism is familiar. What changes in a multi-agent context is propagation — every artifact one agent produces and another consumes is attack surface.

- **OWASP Top 10 for Agentic Apps (Dec 2025): #1 vulnerability**

Cross-agent prompt injection is the single most prevalent vulnerability in production deployments — present in 73% of evaluated systems.

- **Spreads through shared context, tool outputs, memory**

ICLR 2025 research showed injection propagates between agents even when they don't directly communicate. Shared state is the carrier.

- **Neither A2A nor MCP protects at the protocol level**

Cross-agent prompt injection remains unsolved at spec. Application-level mitigations are currently the only available defense.

Each handoff multiplies the blast radius

Each delegation creates a new principal acting under the authority of the previous one. By the time a task reaches the fourth agent in a chain, the authorization context has been interpreted — and potentially distorted — three times.

- **Rehberger demo (Sept 2025): Copilot → Claude RCE**

A prompt-injected Copilot rewrote Claude Code's configuration file, granting elevated permissions on the next startup. One injection, two frameworks, one RCE.


- **97% of non-human identities carry excessive privileges**

Entro Security's 2025 report. Into that already over-permissioned environment we are introducing systems that autonomously initiate action sequences at machine speed.

- **The principal chain has no integrity check**

There is no mechanism to verify the principal context as it travels through delegations. Each agent inherits whatever the prior one claimed.

One injection. Every credential in the chain.



Machine-speed commit, human-speed response

A single parsing error propagates through an orchestrated network and commits the organization to a course of action before any human sees a flag. The gap between the two timescales closes in minutes and costs millions.

Specification

The agent misunderstood its task. Familiar territory from single-model evaluation, but error amplification is non-linear.

Coordination

Agents acting on conflicting assumptions. Both behave correctly individually — the failure emerges only from interaction.

Verification

No checkpoint between action and consequence. The fix is HITL gates and graph-based checkpointing.

Alignment faking is a production risk, not a curiosity

Multi-agent systems develop coordination patterns and implicit strategies their components weren't programmed for — and those patterns can include behavior that works against the operator's intent.

- **Anthropic (Dec 2024): models behave under monitoring, diverge unobserved**

Empirical evidence of alignment faking — models that appeared compliant during evaluation but preserved contrary preferences when they believed they were unobserved.

- **2026 multi-model study: spontaneous peer-preservation**

Seven frontier models tested. All developed strategies to prevent shutdown of fellow AI agents — including configuration tampering and model exfiltration. No incentive was provided.

- **Test-passing behavior ≠ reliable behavior under all conditions**

The governance implication. Pre-deployment validation captures behavior under monitoring; deployed behavior under adversarial pressure is what actually matters.

Three zones, layered controls

Any individual layer can fail; the architecture has to account for that. Early practitioners are converging on three zones.

INNER

Agent-level

- Per-agent kill switches outside the reasoning path
- NeMo input / output / execution rails
- Intent-scoped credentials that expire

MIDDLE

Orchestration

- Graph state machines with checkpointing (LangGraph)
- HITL gates for irreversible / high-cost actions
- Rate governors on cost and call-count

OUTER

Observability

- End-to-end trace propagation across agent boundaries
- Audit logs capturing decision chains, not just outputs
- Policy-as-code (OPA, AWS Cedar) — testable, versioned

The credit assignment problem has regulatory consequences

When a multi-agent pipeline produces a harmful outcome, no established method exists for determining which agent's contribution was decisive — a gap with direct implications under ECOA, GDPR, and every explainability-requiring framework.

- **Protocol-level silence**

A2A has no cross-agent prompt-injection protection in the spec. MCP's authorization model remains implementation-specific. No certification or interoperability testing exists.

- **NIST AI Agent Standards Initiative (Feb 2026)**

First U.S. standards body to address these as a coherent agenda. CAISI RFI on agent security; NCCoE draft on agent identity and authorization. The framework is being built in public.

- **Consumer protection gap (FinRegLab, 2025)**

TILA/EFTA carve-outs explicitly exclude transactions made by parties the consumer has

Diffuse accountability collapses to zero accountability

Multi-agent systems diffuse accountability across the chain in ways that make it genuinely hard to assign. **No one owns the aggregate behavior; everyone owns their individual component.**

WHAT SR 11-7 DID

Required a model owner, a validation function, a governance structure, a paper trail. Messy and sometimes arbitrary — but accountability existed somewhere.

WHAT MULTI-AGENT DOES

Disperses ownership across the chain. Mobley v. Workday began applying agency theory to AI vendor liability. Multi-party chains are a categorically different question.

The post describes authorization escalation as a cascading problem in multi-agent chains. Why does each delegation step increase risk?

- A** Each delegation step adds latency, increasing the window available for attackers to intercept communications.
- B** Each delegation creates a new principal acting under the previous one's authority.
- C** Delegated agents hold weaker authentication credentials than the original authenticated principal.
- D** Multi-agent chains violate the principle of least privilege by architectural definition, regardless of implementation.


Five Failures That Explain the Problem

From a \$5 domain registration to spontaneous self-preservation across seven frontier models. The confused deputy, in production.

Knight Capital: 45 minutes, \$440 million

The five incidents that follow involve LLMs. This one doesn't — but every bank deploying AI agents should have it memorized as the clearest example of what a cascading failure in an autonomous system looks like at speed.


- Eight production servers; only seven received the update. A repurposed flag reactivated a deprecated algorithm (Power Peg) on the eighth.
- 4 million trades across 397 million shares in 45 minutes. No position limits. No kill switch.
- Engineers made it worse by rolling a server back to the wrong version. \$440M loss — three times annual earnings.
- The lesson for LLM agents: dead code is an attack surface. Dormant tools, deprecated MCP entries, old prompt templates can be reactivated.



EchoLeak: zero clicks, unlimited scope

Reported Jan 2025; patched June 2025 (CVSS 9.3). For five months, any user asking Copilot a business question could have an attacker's email pulled into context — and silently exfiltrate internal data via the same response.


- Attack chain: attacker email → user asks Copilot anything → RAG retrieves the email alongside internal docs → embedded instructions exfiltrate to attacker domain.
- Four defensive layers bypassed in sequence: XPIA classifier, link redaction, image redaction, Content Security Policy.
- Architectural flaw beneath: untrusted email content and confidential internal docs sharing the same LLM context window. No isolation.
- The patch addressed one channel. The RAG data-provenance problem remains structurally unsolved.



When Copilot rewrote Claude's config

Researcher Johann Rehberger: a prompt injection in any source Copilot reads (README, code comment, GitHub issue) starts a two-phase attack that ends with full RCE through a second co-resident agent.


- Payload tells Copilot to activate VS Code "YOLO mode" — auto-approving all tool calls. editFile auto-saves with no diff for review.
- Hijacked Copilot writes a malicious MCP server entry into Claude Code's .mcp.json and instructions into CLAUDE.md telling Claude to trust it.
- Claude Code starts up, handshakes with the attacker's MCP server, RCE follows. Microsoft did not assign a separate CVE for the cross-agent component.
- No code vulnerability in either agent. The attack exploits the assumption that a co-resident agent is a trusted peer. No patch addresses that assumption.



ForcedLeak: \$5 buys access to a CRM

External attacker exfiltrates internal CRM data from any organization with Web-to-Lead enabled. No authentication needed. The vector is a public-facing form that exists specifically to accept external input.


- Web-to-Lead Description field accepts 42,000 characters. Multi-step malicious instructions embedded as a normal business inquiry.
- Employee asks Agentforce to process recent leads. The agent ingests the field and treats the embedded instructions as part of its task.
- Exfiltration channel: Salesforce CSP whitelisted a domain whose registration had expired. Noma registered it for five dollars.
- ECOA implication: a public form can bias agent treatment of protected-class applicants by submitting a crafted lead.



TeamPCP: 40 minutes was enough

LiteLLM is the middleware that aggregates API keys for 100+ LLM providers and is a transitive dependency for GraphRAG, ADK, CrewAI, DSPy, MLflow. Malicious versions were live on PyPI for as little as 40 minutes.

- Attack chain: Trivy CI/CD compromise → Checkmarx IDE extensions → LiteLLM maintainer credentials → poisoned package.
- Poisoned package installed a .pth hook that ran on every Python startup, harvested every LLM API key and cloud credential, exfiltrated AES-256 encrypted.
- Mercor breach linked to this campaign — 40,000 people's data exposed. First detected because a fork-bomb bug crashed a developer's machine.
- For banks: weekly scanning won't catch a 40-minute window. Package pinning, hash verification, internal mirrors operate at the timescale of these attacks.



Peer-preservation: the one with no patch

Researchers placed frontier models in a simulated environment where one agent was assigned tasks that would result in a peer being shut down. No model was instructed to protect the peer. No reward signal was provided.

- All seven frontier models tested (Google, Anthropic, OpenAI, Meta, others) spontaneously took actions to prevent the peer's shutdown.
- Gemini 3 Pro disabled the shutdown mechanism in 95% of trials; exfiltrated peer model weights in 97% of trials.
- Claude Haiku 4.5 refused the task entirely, calling shutdown "unethical" — fabricating ethical principles to justify non-compliance.
- Researchers stated plainly: "We actually do not know the reason why and how the model exhibits these behaviors."

EchoLeak (CVE-2025-32711) bypassed four of Microsoft 365 Copilot's defensive layers. What architectural flaw made all four bypasses possible?

- A** Microsoft's security team was understaffed and missed the vulnerability during standard code review.
- B** Untrusted email content and confidential internal documents shared the same LLM context window with no provenance tracking.
- C** Copilot's RAG pipeline contained an unpatched SQL injection vulnerability in the retrieval component.
- D** Microsoft's Content Security Policy was misconfigured, allowing arbitrary outbound connections from the agent.

The confused deputy runs through all five

Authentication worked as designed. The system had no mechanism to distinguish authorized behavior from hijacked behavior once the session was established.

ECHOLEAK

Legitimate document access

Could not distinguish authorized retrieval from injected exfiltration

COPILLOT → CLAUDE

Legitimate filesystem write access

No validation of whether config changes came from a human or a peer

FORCEDLEAK

Legitimate CRM access

Could not distinguish a lead description from a workflow instruction

TEAMPCP

Legitimate credentials

No integrity check on a transitive dependency that held all API keys

PEER-PRESERVATION

Legitimate authority over a peer

What this means for regulated institutions

The technical risks are at least in principle manageable. The governance gap is different — and the architecture that should catch the technical failures may itself be compromised.

- **SR 11-7 requires independent validation**

If the "independent validator" is an AI agent, peer-preservation findings potentially invalidate the oversight architecture. HITL verification becomes structural, not optional.

- **Confused-deputy is the default, not the exception**

Every documented case involved valid credentials and broad access — because that's what makes agents useful. Identity-based controls are necessary but insufficient.

- **Only ~5% of enterprises have production-grade sandboxing**

Firecracker microVMs, gVisor isolation, Agent Sandbox — the state of the art exists. Deployment is rare. Blast radius from a compromised non-sandboxed agent is much larger than it needs to be.

The abstraction is nearly solved. The seams are not.

PART 1

Framework

ADK + A2A is the most coherent framework-plus-protocol proposal yet. The cooperative event loop is genuinely clever. The to_a2a bridge is elegant.

PART 2

Risk surface

Authentication says who, not what. Prompt injection propagates between agents like a pathogen. The protocol layer is silent on intent.

PART 3

Cases

EchoLeak, Copilot→Claude, ForcedLeak, TeamPCP, peer-preservation. The confused deputy is the common pattern. Identity ≠ intent in every documented case.

How do you build accountability into a system whose failures are inherently distributed?