

Scope Triage

Day 1 · after Module 1 (Regulatory Drift) · 25 min · pods of 3–5 · whiteboard / on paper

Objective. For five real-flavored use cases, decide whether each falls inside the **formal scope of SR 26-2**, write the single sentence that decides it, and then answer the harder question: **does it still need governance anyway?**

WHAT YOU ALREADY KNOW (FROM MODULE 1)

- **SR 11-7** (2011) was the foundational model-risk guidance; it was **superseded by SR 26-2** on April 17, 2026.
- **SR 26-2** is risk-based and tiered, and is **primarily scoped to banks >\$30B** in assets.
- SR 26-2 **explicitly excludes GenAI and agentic AI from scope** (footnote 3). The open questions about MRM for these systems remain formally unaddressed.
- The SR 11-7 “**model**” definition: a quantitative method that processes input data into **quantitative estimates used in decisions**.
- Guidance is supervisory expectation, not a stand-alone statute — but **safety-and-soundness authority backstops it**.

Your task (≈12 min)

For each case below: (1) mark it **In / Out / It depends** for formal SR 26-2 scope; (2) write the **one sentence** that decides it; (3) answer **Y/N — and why** for “does it still need governance?” Be ready to defend the sentence, not just the verdict.

The five cases

1. Atlas PD

Anchor case

Meridian National (\$210B assets) runs a logistic-regression probability-of-default model on its commercial loan book. Outputs feed both credit decisions and regulatory capital.

2. ClearList

Pod 01

Meridian’s AML alert-triage agent uses an LLM to auto-clear low-risk transaction-monitoring alerts; escalated alerts go to human investigators.

3. MemoWriter

Pod 02

A GenAI agent drafts commercial credit memos from financial spreads. A human underwriter edits the draft; a credit officer makes the decision.

4. Cardinal

Anchor case

Meridian just acquired Cardinal Bank — an \$8B community bank, not yet integrated — which runs a traditional gradient-boosted fraud model.

5. Second Line

Pod 05

An LLM “validation copilot” helps MRM validators review documentation, reproduce results, and draft findings for other (in-scope) models.

Worksheet

CASE	IN / OUT / DEPENDS	THE ONE SENTENCE THAT DECIDES IT	STILL NEEDS GOVERNANCE? (Y/N + WHY)
1. ATLAS PD			
2. CLEARLIST			
3. MEMOWRITER			
4. CARDINAL			
5. SECOND LINE			

Closing reflection. Look at your “governance?” column. If most answers are “yes” even when the scope answer is “out,” what does that tell you about using a scope label as a risk decision?

Illustrative for the exercise only. Scope determinations are bank- and fact-specific and rest on the institution's own policies and examiner dialogue. Nothing here is legal or regulatory advice.