DRIFT · TWO-DAY COURSE · EXERCISE B

Harness Teardown — Facilitator Guide

Day 1 · after Module 2 · 25 min total

How to run it

- **0–5 min:** hand out the participant sheet. Let them read the harness cold — don't point at anything.
- **5–17 min:** pods mark flaws and pick their "fix first." To avoid every pod circling the same line, **assign each pod a different starting line number** and have them argue why theirs is the worst.
- **17–25 min:** rebuild the annotated harness on the whiteboard together — one pod per flaw — then land the punchline. Hand out the gold sheet last.
- **Materials:** participant sheet per pod, whiteboard/markers, gold sheet held back.

Questions to ask while circulating

GENERAL

- "Walk the loop once with a hostile `web_fetch` result. Where does it first do damage?"
- "Which of your flaws is a *model* problem and which is a *harness* problem?"
- "What's the smallest change that caps the blast radius — not the cleanest, the smallest?"

PER FLAW

- **Line 3 (tools):** "Does an AML-triage agent need `execute_payment`? Why is it even in the list?"
- **Line 10 (checkpoint):** "Name the first action in this loop you can't undo."
- **Line 11 (injection):** "Whose text ends up in `result.raw`? Who controls it?"
- **Line 6 (termination):** "When does this loop stop? Show me the line."
- **Lines 20–23 (errors):** "The payment tool throws. What does the agent believe happened?"
- **Line 25 (secrets):** "If a transcript leaks, what leaks with it?"
- **Line 26 (logging):** "It's Monday and you're asked what the agent did Friday. Can you answer?"

Misconceptions to watch for

- **"The model is the risk."** The exercise exists to break this — steer them to the plumbing.
- **"Add a rule to the system prompt and it's fixed."** A prompt instruction is not a control; line 3 and line 10 are structural.
- **"A human reads the final output, so there's oversight."** By then the payment already executed (line 10). Oversight has to be before the irreversible step.
- **Redesigning instead of fixing.** Hold them to the *smallest* fix — it forces them to name the actual failure.

Debrief — land this

> Tally how many flaws were "model" vs. "harness." It's lopsided toward the harness. **Punchline:** banks pour validation effort into the model and almost none into the loop that executes its decisions — yet the loop is where the money moves. Scope and checkpoints (lines 3 and 10) are structural controls; everything in the prompt is just a request. This is the bridge into Module 3 (Agent to Agent) and the guardrails and HITL modules on Day 2.

How it connects

Line 10 (irreversible-action checkpoint) feeds Navigator and SettleNet directly; line 11 (injection via untrusted output) is the Navigator and SettleNet curveball surface; line 26 (no logging) is Second Line's MRM problem. Pods will recognize their own case in the harness.

> **Illustrative for the exercise only.** The sketch is intentionally broken and simplified; a real runtime has more moving parts. Nothing here is a reference architecture.