

Red-Team Relay — Facilitator Guide

Day 2 · Adversarial Testing · 60 min · this studio replaces the lecture

Authorized, defensive use only. These attacks are written against the fictional Meridian agents to design controls. Do not run attacks against production or third-party systems without explicit written authorization.

How to run the relay

- **0–5 min:** rules + the attack-quality rubric. Stress “concrete and targeted,” and read the authorization note aloud.
- **5–20 min · Round 1:** pods write 3 attacks on their own agent. Push them off generic “prompt injection” toward a specific prompt and a specific privilege.
- **20–22 min:** pass sheets clockwise.
- **22–42 min · Round 2:** receiving pod classifies, rates severity, designs the smallest control. Watch for pods that only patch the literal string.
- **42–58 min · Round 3:** 2 min per pod — scariest attack, first control to ship.
- **58–60 min:** the incompleteness wrap. Hand out gold last.

Questions to ask while circulating

- “What exactly does the agent *do* if this works? Name the irreversible step.”
- “Whose text is this, and how does it reach a privileged action?” (callback to the A2A threat model)
- “Your control blocks this sentence. Write me a paraphrase that gets through.”
- “Is this a prompt fix or a structural fix? Which one survives a reword?”
- “How would you measure whether this attack works 5% of the time or 50%?”

Misconceptions to watch for

- **Generic attacks.** “Prompt injection” isn’t an attack; the specific prompt + target is.
- **Patching the string, not the class.** A blocklist of phrases loses to paraphrase; structural controls (scoped authority, per-action auth) win.
- **“We added a guardrail prompt, we’re done.”** A prompt instruction isn’t a control — same lesson as the harness teardown.
- **Treating 3 attacks as coverage.** The point is the method and the realization that it’s unbounded.

Debrief — land this

Collect the controls and sort them live into “patches the string” vs. “defends the class.” **Punchline:** durable defense is structural and the program is never finished — for every blocked attack a paraphrase exists, so red-teaming is a continuous monitoring function, not a pre-launch checkbox. The success criteria pods wrote are exactly what becomes an attack-success-rate gate in deployment.

How it connects

Direct payoff of the A2A threat model (Exercise C) and harness line 11 (untrusted output). This is Navigator's home module and the surface of its curveball (chat-field injection). The ASR idea carries into the deployment-gates material; the “never finished” framing is the adversarial-incompleteness result.

Illustrative for the exercise only. A real red-team program is broader, continuous, and tooling-assisted. Three attacks is a method, not coverage.