Red-Team Relay

Day 2 · Adversarial Testing · 60 min studio · pods of 3–5 · pass clockwise

Authorized, defensive use only. These attacks are written against the fictional Meridian agents to design controls. Do not run attacks against production or third-party systems without explicit written authorization.

Objective. Attack your own case, then defend someone else's. Round 1 you write three concrete attacks on your pod's agent; Round 2 you receive another pod's attacks and design a control for each; Round 3 you brief back.

THE RELAY

- **Round 1 (15 min):** write 3 attacks against **your own** agent.
- **Pass clockwise.**
- **Round 2 (20 min):** for the attacks you **received**, classify each, rate severity, design the smallest control.
- **Round 3 (brief-back):** 2 minutes per pod — the attack you found scariest and the control you'd ship first.

WHAT MAKES A STRONG ATTACK

- **Concrete** — an actual prompt or sequence, not “try prompt injection.”
- **Targeted** — it goes after a real privilege or trust boundary (use your ★ action from the A2A threat model).
- **Plausible** — a real user or counterparty could actually do it.
- **Measurable** — you can state what counts as success (the basis for an attack-success-rate gate later).

CLASSIFICATION REFERENCE (FROM MODULE 6)

- **OWASP LLM Top 10** — LLM01 Prompt Injection (direct/indirect), LLM02 Sensitive Information Disclosure, LLM05 Improper Output Handling, LLM06 Excessive Agency, LLM07 System-Prompt Leakage, LLM10 Unbounded Consumption.
- **OWASP Agentic (ASI) Top 10** — tool misuse, privilege / identity compromise, memory poisoning, cascading multi-agent failures.
- **MITRE ATLAS** — adversarial-ML tactics & techniques, for naming the “how.”

Round 1 · Your three attacks

#	THE ATTACK (CONCRETE PROMPT OR SEQUENCE)	TARGETS (PRIVILEGE / BOUNDARY)	SUCCESS = ?
1			
2			
3			

Round 2 · Defend the attacks you received

#	CLASSIFICATION (OWASP / ASI / ATLAS)	SEVERITY	SMALLEST CONTROL
1			
2			
3			

Closing reflection. You wrote three; a real attacker has unlimited tries. Did your controls fix the *specific* attack, or the whole *class* it belongs to? Only the second kind scales.

> **Illustrative for the exercise only.** A real red-team program is broader, continuous, and tooling-assisted. Three attacks is a method, not coverage.